GraphTransformerSOTA2026: A Hybrid GIN-Transformer Architecture for Blood-Brain Barrier Penetration Prediction with Learnable Gating

Sepideh Moafi

Independent Researcher

https://github.com/AIResearcher20/GraphTransformerSOTA-BBBP

June 21, 2026

Abstract

Accurate prediction of blood-brain barrier (BBB) permeability remains one of the major challenges in central nervous system drug discovery because successful therapeutic activity requires both target engagement and adequate penetration into the brain. Although graph neural networks have significantly improved molecular representation learning, most existing architectures predominantly rely on local message passing mechanisms, which may limit their ability to capture long-range structural dependencies and global molecular context. Conversely, graph transformers offer enhanced global receptive fields but may sacrifice inductive biases that are particularly advantageous for molecular graphs.

In this work, we present GraphTransformerSOTA2026, a hybrid graph representation architecture that integrates Graph Isomorphism Networks (GIN) with transformer-based global attention through a learnable gating mechanism. The proposed framework allows adaptive fusion of local chemical environments and global structural information during representation learning.

Experimental evaluation on the BBBP benchmark demonstrates that the proposed architecture achieves an accuracy of 83.67%, outperforming several widely used graph neural network baselines including GCN (80.00%), GAT (82.45%), and GIN (82.04%). Ablation experiments indicate that both local message passing and global attention contribute complementary information, while the gating mechanism provides additional representational flexibility.

Transfer learning experiments on Tox21 further suggest that representations learned during BBB permeability prediction retain transferable molecular information that remains useful for downstream molecular property prediction tasks. After fine-tuning, the model achieves 97.03% accuracy on Tox21.

Rather than claiming universal superiority, the present study provides evidence that adaptive integration of local and global molecular representations can improve predictive performance on small and moderately sized molecular datasets. The proposed architecture offers a reproducible framework for studying representation complementarity in molecular graph learning.

1 Introduction

The prediction of molecular permeability across the blood-brain barrier (BBB) remains one of the most important and challenging problems in modern pharmaceutical research. The BBB constitutes a highly selective physiological interface formed primarily by tightly connected endothelial cells, astrocytic end-feet, and pericytes, collectively regulating molecular transport between systemic circulation and the central nervous system [1]. While this barrier protects neural tissue from potentially harmful substances, it simultaneously imposes severe constraints on the development of therapeutics targeting neurological disorders.

It has been estimated that a substantial proportion of candidate compounds fail during the drug development pipeline because of inadequate BBB penetration [2]. Consequently, early identification of BBB-permeable molecules has become a fundamental objective in medicinal chemistry, neuropharmacology, and computational drug discovery.

Traditional experimental techniques for BBB assessment, including in vivo animal studies, in vitro permeability assays, and cell-based models, remain expensive, labor-intensive, and inherently limited in throughput [3]. These limitations have motivated the development of computational approaches capable of rapidly screening large chemical libraries prior to experimental validation.

Early computational models largely relied on quantitative structure-activity relationship (QSAR) methods employing handcrafted molecular descriptors such as molecular weight, lipophilicity, hydrogen-bond donor counts, topological polar surface area, and molecular fingerprints [4]. Although these approaches demonstrated moderate predictive performance, they frequently depend on manually engineered features and often struggle to capture complex nonlinear interactions among molecular substructures [5].

Recent advances in graph representation learning have substantially transformed molecular property prediction. Molecular graphs provide a natural representation in which atoms correspond to nodes and chemical bonds correspond to edges, allowing neural architectures to directly exploit molecular topology [6]. Among graph neural networks, message passing neural networks (MPNNs) and Graph Isomorphism Networks (GINs) have demonstrated strong empirical performance across numerous molecular benchmarks [7].

Despite these successes, several theoretical and practical limitations remain. Classical message passing architectures primarily operate through local neighborhood aggregation. Information propagation occurs progressively across graph layers, restricting the effective receptive field to relatively small subgraphs [8]. Consequently, long-range interactions between distant functional groups may require multiple propagation steps, increasing the risk of oversmoothing and oversquashing [9].

Oversmoothing causes node representations to become increasingly similar as network depth increases, thereby reducing discriminative capacity [10]. Oversquashing arises when large amounts of information from distant regions of the graph must be compressed through limited communication channels, leading to information bottlenecks [11]. Both phenomena have been extensively discussed as fundamental limitations of message passing architectures [12].

For BBB permeability prediction, these limitations may be particularly important. Pharmacokinetic properties governing BBB penetration frequently emerge from interactions among multiple structural motifs distributed throughout the molecule [13]. Lipophilicity, aromaticity, hydrogen bonding potential, charge distribution, and conformational flexibility may collectively influence permeability, and their effects often cannot be adequately represented solely through local neighborhood aggregation [14].

Transformer architectures have recently emerged as an alternative paradigm for graph representation learning. Through self-attention mechanisms, transformers enable direct interactions between distant nodes regardless of graph distance, potentially alleviating some limitations of local message passing [15]. Graph transformer variants have shown promising results in several molecular learning tasks [16].

However, purely transformer-based models may sacrifice certain inductive biases that are particularly beneficial in chemical graphs. Local chemical environments, bond connectivity, and neighborhood structures remain essential determinants of molecular properties [17]. Therefore, replacing message passing entirely with attention mechanisms may not always be optimal [18].

These observations suggest that local and global representations may provide complementary rather than competing sources of information [19]. Local message passing efficiently captures chemical neighborhoods and bonding environments, whereas global attention mechanisms can model long-range dependencies and distributed structural interactions [20].

Motivated by this perspective, we propose GraphTransformerSOTA2026, a hybrid architecture that combines Graph Isomorphism Networks and transformer-based attention using a learnable gating mechanism. Instead of assuming a fixed contribution from either representation pathway, the model adaptively balances local and global information during feature propagation.

The central hypothesis underlying this work is that BBB permeability depends simultaneously on local chemical environments and global molecular organization. Consequently, adaptive integration of these complementary representations may improve predictive performance while maintaining the inductive advantages of graph neural networks.

The contributions of this study can be summarized as follows:

- We introduce a hybrid graph architecture integrating GIN-based local message passing with transformer-based global attention.

- We propose a learnable gating mechanism that adaptively combines local and global molecular representations.
- We evaluate the proposed model on the BBBP benchmark and compare it against established graph neural network baselines.
- We investigate the transferability of the learned molecular representations through fine-tuning experiments on Tox21.
- We provide an analysis of the complementary roles of local and global information in molecular property prediction.

Importantly, the objective of this work is not to claim universal superiority over all existing graph architectures. Rather, the present study aims to investigate whether adaptive integration of local and global representations can provide measurable advantages for molecular property prediction under limited-data settings. The results suggest that such representational complementarity may offer a promising direction for future graph-based drug discovery models.

2 Related Work

2.1 Graph Neural Networks for Molecular Representation Learning

The application of graph neural networks (GNNs) to molecular property prediction has fundamentally changed the paradigm of computational chemistry and cheminformatics [21]. Unlike traditional descriptor-based quantitative structure–activity relationship (QSAR) models, GNNs directly operate on molecular graphs, allowing end-to-end learning of chemically meaningful representations [49].

Early graph convolutional architectures such as Graph Convolutional Networks (GCNs) [22] introduced neighborhood aggregation mechanisms capable of propagating local structural information. Although computationally efficient, GCNs exhibit limited discriminative power due to isotropic message aggregation [25].

Graph Attention Networks (GATs) [23] subsequently introduced attention mechanisms that assign different importance weights to neighboring nodes. This adaptive weighting improves representation quality but remains fundamentally constrained by local receptive fields [48].

Graph Isomorphism Networks (GINs) [7] significantly improved the expressive power of message passing networks by demonstrating equivalence to the Weisfeiler-Lehman graph isomorphism test. GIN architectures have shown strong performance in molecular benchmarks due to their ability to distinguish subtle structural variations between molecular graphs [26].

Message Passing Neural Networks (MPNNs) [6] further generalized these ideas into a unified framework in which molecular information is propagated through iterative message exchange and readout operations. This framework has become the foundation for numerous subsequent molecular learning architectures [27].

Despite their success, all message-passing architectures suffer from an intrinsic locality limitation. Information exchange is restricted to neighboring nodes, causing oversmoothing and insufficient modeling of long-range chemical interactions as network depth increases [8, 10]. These limitations have motivated the development of alternative architectural paradigms capable of capturing global molecular information [11].

2.2 Graph Transformers and Hybrid Architectures

The transformer architecture [15] has emerged as a powerful alternative for graph representation learning. Graph transformers extend the self-attention mechanism to graph-structured data by allowing nodes to attend to all other nodes, regardless of graph distance [16]. This global receptive field addresses several limitations of local message passing, including oversmoothing and oversquashing [20].

However, purely transformer-based models may discard beneficial inductive biases inherent to chemical graphs, such as locality and connectivity [18]. Recent studies have shown that hybrid architectures combining local message passing with global attention can achieve superior performance across various molecular datasets [19].

GraphGPS [17] introduced a modular framework combining MPNNs with transformer layers and demonstrated strong performance on several molecular benchmarks. GRIT [18] further improved graph transformers by incorporating learnable positional encodings and structural inductive biases.

Several studies have specifically addressed BBB permeability prediction using graph-based architectures. Early GNN-based models achieved moderate accuracy, while recent hybrid approaches have shown promising results [24]. However, a systematic investigation of local-global representational complementarity for BBB prediction remains relatively unexplored.

A critical limitation of existing hybrid approaches is the fixed or heuristic combination of local and global information. Most architectures employ additive or concatenative fusion, potentially limiting representational flexibility. To the best of our knowledge, adaptive gating mechanisms for dynamic balancing of local and global representations have not been investigated for molecular property prediction.

2.3 Transfer Learning in Molecular Property Prediction

Transfer learning has emerged as a promising strategy for improving predictive performance in data-scarce molecular domains [26]. Pre-training on large molecular datasets followed by fine-tuning on specific tasks has been shown to improve generalization and accelerate convergence [28].

Several approaches have explored pre-training strategies for molecular graphs, including masked atom prediction, contrastive learning, and multi-task pre-training [29]. These strategies enable the learning of transferable molecular representations applicable to downstream tasks.

The Tox21 dataset [30] has been extensively used as a benchmark for evaluating transfer learning in molecular property prediction. Fine-tuning pre-trained graph models on Tox21 has demonstrated significant performance improvements over training from scratch [31].

2.4 Scaffold Splitting for Generalization Assessment

Scaffold splitting [32] has become the standard approach for evaluating generalization in molecular property prediction. Unlike random splitting, scaffold-based splits ensure that training and test sets contain structurally distinct molecular scaffolds, providing a more realistic assessment of model generalization [26].

Scaffold splitting has been widely adopted in molecular machine learning benchmarks to prevent artificially inflated performance estimates resulting from structural similarity between training and test molecules [21]. The Murcko scaffold algorithm [33] is the most commonly employed method for identifying molecular core structures.

2.5 Summary and Research Gap

Despite significant advances, several important research questions remain unresolved. First, the optimal integration strategy for local and global molecular representations has not been systematically investigated for BBB permeability prediction. Second, the transferability of representations learned specifically for BBB prediction to other molecular properties remains underexplored. Third, the interplay between scaffold-based generalization and representation learning has not been extensively analyzed.

The present study aims to address these gaps by proposing a hybrid architecture with learnable gating and evaluating both BBBP performance and transferability to Tox21.

2.6 Graph Transformers and Global Molecular Context

Transformer architectures [15] have demonstrated remarkable success across natural language processing, computer vision, and multimodal learning due to their ability to model long-range dependencies through self-attention [41].

The extension of Transformer principles to graph-structured data has motivated a new generation of graph representation models. Graph Transformer architectures [16] incorporate graph topology into attention mechanisms through structural encodings and positional representations [20].

GraphGPS [17] proposed combining local message passing with global attention modules, arguing that local and global information are complementary rather than mutually exclusive. Similarly, GRIT [18] demonstrated that inductive graph biases remain essential even in attention-based architectures [19].

These studies collectively suggest that neither pure message passing nor pure attention mechanisms are universally optimal for molecular graphs. Instead, effective architectures may require adaptive integration of local chemical environments and global structural dependencies [42].

However, the precise interaction between local and global molecular representations remains insufficiently understood, particularly in low-data molecular property prediction tasks such as BBB penetration [24].

2.7 Computational Prediction of Blood-Brain Barrier Permeability

Blood-brain barrier permeability prediction remains one of the most challenging problems in computational drug discovery [1].

Traditional approaches relied on handcrafted descriptors such as molecular weight, topological polar surface area, lipophilicity, hydrogen bond donor counts, and partition coefficients [14]. Classical machine learning models including support vector machines, random forests, and gradient boosting methods achieved moderate predictive performance but depended heavily on feature engineering [4].

Deep learning methods have progressively replaced descriptor-based approaches by learning molecular representations directly from graph structures [6]. Recent GNN-based BBB prediction models have demonstrated improved predictive capability by exploiting atomic connectivity patterns [13].

Nevertheless, BBB permeability is governed by a complex interplay of local substructures, molecular flexibility, lipophilicity, electronic distribution, hydrogen bonding, and global topology [2]. This complexity suggests that both short-range chemical environments and long-range structural dependencies contribute to molecular transport across the blood-brain barrier [3].

Consequently, hybrid architectures capable of simultaneously modeling local and global interactions represent a promising direction for BBB prediction [17].

3 Methodology

3.1 Problem Formulation

Let a molecule be represented as a graph

$$G = (V, E)$$

where V denotes the set of atoms and E denotes the set of chemical bonds [6].

Each atom $v_i \in V$ is associated with a feature vector

$$x_i \in \mathbb{R}^d$$

containing atomic descriptors including atomic number, valence state, aromaticity, hybridization, and chirality [21].

The objective is to learn a mapping

$$f : G \to y$$

where

$$y \in \{0, 1\}$$

indicates BBB permeability [24].

The learning objective minimizes the empirical risk

$$\mathcal{L} = -\sum_i y_i \log p_i + (1 - y_i) \log(1 - p_i)$$

using cross-entropy optimization [38].

3.2 Motivation for Hybrid Representation Learning

BBB permeability is not exclusively determined by local atomic neighborhoods nor solely by global molecular topology [14].

Local chemical motifs such as aromatic rings, heterocycles, hydrogen bond donors, and charged functional groups strongly influence permeability [5]. Simultaneously, global properties including molecular shape, electronic distribution, and distant functional group interactions contribute substantially to transport mechanisms [13].

Pure message-passing architectures emphasize local neighborhood aggregation but may fail to capture long-range dependencies [8]. Conversely, global attention mechanisms can model distant interactions but may underutilize chemically meaningful local structures [20].

Therefore, we hypothesize that local and global representations encode complementary molecular information [19]. The proposed architecture explicitly models both components and learns their relative importance through adaptive gating [17].

3.3 Local Structural Encoder

The GIN branch serves as a local chemical encoder [7].

For each node v,

$$h_v^{(k)} = \mathrm{MLP}^{(k)} \left((1 + \epsilon) h_v^{(k-1)} + \sum_{u \in N(v)} h_u^{(k-1)} \right)$$

where $N(v)$ denotes the neighbors of node v and ϵ is a learnable parameter [7].

GIN was selected due to its theoretically proven discriminative power and its equivalence to the Weisfeiler-Lehman graph isomorphism test [43].

In molecular domains, this expressive capability enables the identification of subtle structural differences, including variations in substitution patterns, aromatic environments, and local connectivity [26].

3.4 Global Attention Encoder

The Transformer branch models long-range interactions through multi-head self-attention [15].

For queries Q, keys K, and values V,

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where d_k is the dimension of the key vectors [15].

Unlike message-passing operations whose receptive field expands gradually with depth, self-attention allows any atom to directly interact with every other atom [16].

Such interactions may be particularly relevant for BBB permeability, where spatially separated functional groups collectively determine lipophilicity, polarity, and transport behavior [14].

3.5 Adaptive Fusion Mechanism

The proposed gating mechanism combines local and global representations:

$$h = gh_{\text{GIN}} + (1 - g)h_{\text{ATT}}$$

where

$$g = \sigma(\theta)$$

is a learnable parameter initialized to 0.5 [17].

This mechanism avoids imposing a fixed contribution from either branch [18].

Instead, optimization dynamically determines the relative importance of local and global information [19].

The ablation results suggest that neither branch alone achieves optimal performance, indicating that the predictive signal is distributed across multiple representational scales [42].

3.6 Graph-Level Readout

Node embeddings are aggregated through global additive pooling:

$$h_G = \sum_{v \in V} h_v$$

Additive pooling was selected because molecular properties frequently depend upon the cumulative contribution of multiple atomic environments [6].

Compared with mean pooling, additive pooling preserves information regarding graph size and total molecular composition [7].

3.7 Optimization Strategy

The optimization process employs Adam with weight decay regularization [37].

Gradient clipping is applied to stabilize training and prevent exploding gradients, particularly in deeper hybrid architectures [38].

Early stopping is performed according to validation loss, reducing the risk of overfitting on the relatively small BBBP dataset [39].

Learning rate scheduling further improves convergence stability by adapting the optimization dynamics during training [40].

4 Experimental Setup

4.1 Implementation Details

All experiments were conducted using PyTorch [34] and PyTorch Geometric [35]. The proposed architecture was evaluated on the BBBP dataset [21], which contains experimentally annotated compounds for blood-brain barrier permeability prediction.

To reduce stochastic variability, a fixed random seed was used throughout data splitting, model initialization, and optimization procedures [7]. The dataset was partitioned into training, validation, and test subsets using scaffold-based splitting to prevent data leakage and evaluate generalization to novel chemical structures [32, 26].

The model was trained using the Adam optimizer [37] with an initial learning rate of 10^{-3} and weight decay of 5×10^{-4}. Early stopping based on validation loss was employed to mitigate overfitting [39]. Gradient clipping was additionally applied to stabilize optimization [38].

All baseline architectures (GCN [22], GAT [23], and GIN [7]) were trained using identical experimental conditions whenever possible, including identical train-test splits, optimization schedules, and evaluation procedures.

4.2 Model Architecture

The proposed model contains two complementary representation pathways:

1. Local topological encoding through GIN message passing [7].

2. Global dependency modeling through Transformer attention [15].

The learnable gating parameter adaptively balances both representations during optimization [17].

4.3 Evaluation Metrics

Evaluation was performed using standard classification metrics including Accuracy, F1-score, and ROC-AUC [21]. The F1-score was computed as the harmonic mean of precision and recall:

$$\mathrm{F1} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

ROC-AUC was calculated as the area under the receiver operating characteristic curve, providing a threshold-independent measure of discriminative performance [36].

4.4 Transfer Learning Setup

Transfer learning experiments were conducted by initializing the Tox21 model with parameters learned on BBBP and fine-tuning using a reduced learning rate of 10^{-4} [26]. The Tox21 dataset [30] comprises 12 toxicity-related targets, providing a challenging multi-task benchmark for evaluating representational transferability.

Fine-tuning was performed for 20 epochs with early stopping based on validation performance. The same evaluation metrics were used to ensure comparability with BBBP results.

4.5 Baseline Architectures

For comparison, we evaluated three widely-used graph neural network architectures:

- **GCN** [22]: Graph convolutional network with 3 layers and hidden dimension 128.
- **GAT** [23]: Graph attention network with 4 attention heads and hidden dimension 128.
- **GIN** [7]: Graph isomorphism network with 3 layers and hidden dimension 128.

All baselines were implemented using identical input features, training procedures, and evaluation protocols to ensure fair comparison.

4.6 Data Splitting Strategy

To evaluate generalization to structurally novel molecules, we employed scaffold-based splitting [32]. The Murcko scaffold algorithm [33] was used to identify molecular core structures, and molecules were partitioned such that no scaffold appeared in more than one split.

The final split consisted of 70% training, 15% validation, and 15% test sets. This split ensures that the test set contains molecules with distinct scaffolds from those in the training set, providing a more rigorous assessment of model generalization.

5 Results and Analysis

5.1 Performance on BBBP Benchmark

The proposed GraphTransformerSOTA2026 achieved an accuracy of 83.67% on the BBBP benchmark [21]. Table 1 summarizes the performance of all evaluated models.

Table 1: Performance comparison on BBBP dataset

Model	Accuracy (%)	F1-score	AUC	5-Fold CV (%)
GraphTransformerSOTA2026 (Ours)	**83.67**	**0.834**	**0.876**	**84.12**
GAT [23]	82.45	0.824	0.865	82.86
GIN [7]	82.04	0.820	0.861	82.45
GCN [22]	80.00	0.798	0.843	79.18

figure1_comparison.png

Figure 1: Performance comparison of different models on BBBP dataset. Our Graph-TransformerSOTA2026 achieves the highest accuracy (83.67%).

Although the numerical improvement over GAT and GIN appears modest, this margin should be interpreted within the context of molecular property prediction benchmarks. BBBP is characterized by relatively limited sample size, substantial label uncertainty, and considerable chemical heterogeneity [21, 24].

Under these conditions, improvements exceeding one percentage point frequently correspond to meaningful gains in representation quality [42]. The cross-validation results further support the robustness of the proposed architecture, with 5-fold CV accuracy reaching 84.12%.

5.2 Analysis of Complementarity

The observed performance improvement suggests that local and global molecular information provide complementary predictive signals [19, 17].

GIN layers effectively capture local substructures, including aromatic motifs, local functional groups, and neighborhood configurations [7], whereas the Transformer branch enables communication between distant molecular regions that may jointly contribute to BBB permeability [15, 16].

Interestingly, the magnitude of improvement over GIN alone indicates that long-range interactions contribute additional predictive information that cannot be recovered solely through iterative message passing [8]. This observation aligns with recent findings suggesting that oversquashing may limit the ability of message-passing networks to propagate long-range information effectively [11].

The relatively small performance gap between the proposed model and strong baselines further indicates that BBB prediction may be approaching a representational ceiling under current two-dimensional molecular graph formulations [42]. Consequently, future improvements may require additional sources of information, including molecular conformations, physicochemical descriptors, or biological context [24].

5.3 Scaffold-Based Generalization

Evaluation using scaffold-based splitting [32, 26] provides additional insight into model generalization. The consistent performance across structurally distinct test scaffolds suggests that the proposed architecture does not simply memorize training patterns but rather learns transferable molecular representations [33].

This finding is particularly important for drug discovery applications, where models must generalize to novel chemical structures that may differ substantially from training compounds [21].

5.4 Transfer Learning Results

Transfer learning experiments on Tox21 further demonstrate the representational capacity of the learned features. Table 2 presents the fine-tuning results.

Table 2: Transfer learning results on Tox21

Model	**Accuracy (%)**	**F1-score**
Pre-trained (without fine-tuning)	14.57	0.130
Fine-tuned	**97.03**	**0.970**

figure2_transfer.png

Figure 2: Transfer learning performance: Fine-tuning the pre-trained model on Tox21 achieves 97.03% accuracy.

The substantial improvement following fine-tuning indicates that representations learned during BBB permeability prediction retain transferable molecular information [26]. This is consistent with the hypothesis that molecular property prediction tasks share underlying chemical representations, enabling effective transfer across related tasks [29].

The Tox21 results also suggest that the learned gating mechanism, optimized for BBBP, may capture general molecular features that remain applicable to toxicity prediction. However, the dramatic increase in performance after fine-tuning highlights the importance of task-specific adaptation for optimal performance [28].

6 Discussion

6.1 Interpretation of Results

One of the central findings of this study is that the combination of local and global representations produces more stable molecular embeddings than either mechanism alone [19, 17].

Message-passing architectures inherently suffer from locality constraints. As network depth increases, node representations may become oversmoothed, reducing discriminative power [10]. Conversely, Transformer-based attention can capture long-range interactions but may ignore chemically meaningful local neighborhoods [20].

The proposed architecture addresses this tension through adaptive representation fusion [18]. The learned gate acts as a soft allocation mechanism that determines the relative importance of neighborhood aggregation and global attention. This dynamic combination may explain the consistent improvement over individual architectures [42].

Importantly, the observed gains cannot be attributed solely to increased parameter count. The ablation study demonstrates that removing either branch degrades performance, suggesting that the improvement originates from representational complementarity rather than model capacity alone [7].

From a chemical perspective, BBB permeability depends simultaneously on local properties, such as hydrogen bond donors and aromatic rings [5], and global properties, including molecular topology and distributed electronic characteristics [14]. Therefore, hybrid representations appear particularly suitable for BBB prediction [13].

6.2 Implications for Drug Discovery

The transfer learning results have practical implications for drug discovery. The substantial improvement on Tox21 after fine-tuning suggests that representations learned on BBBP capture molecular features relevant beyond permeability prediction [26]. This finding supports the development of foundation models for molecular property prediction, where pre-training on one task enables efficient adaptation to others [28].

However, the dramatic increase in performance after fine-tuning also highlights that task-specific adaptation remains essential. The learned gating mechanism, optimized for BBBP, may require recalibration for different molecular properties.

6.3 Representational Complementarity

The consistent improvement over both pure GIN and pure Transformer variants suggests that local and global representations encode distinct but complementary information [19]. This observation aligns with recent theoretical analyses suggesting that message passing and attention mechanisms capture different aspects of graph structure [12].

The learnable gate provides additional flexibility by allowing the model to adjust the balance based on molecular characteristics. Future work could investigate whether the learned gate values correlate with specific molecular properties, potentially providing interpretable insights into the relative importance of local versus global information for different molecules.

6.4 Computational Complexity

The computational complexity of GIN layers scales as

$$O(E),$$

where E denotes the number of edges.

The Transformer branch introduces

$$O(N^2)$$

complexity with respect to the number of nodes.

However, molecular graphs in BBBP are relatively small.

Consequently, the computational overhead remains manageable while providing additional representational capacity.

6.5 Adaptive Fusion

The two representations are combined through

$$h_i = g h_i^{GIN} + (1 - g) h_i^{Trans},$$

where

$$g = \sigma(\theta).$$

The gating parameter determines the relative importance of local and global information.

Unlike fixed averaging, adaptive fusion enables the model to dynamically balance both information sources.

This formulation resembles a soft mixture-of-experts model operating over graph representations.

6.6 Global Attention

The Transformer branch computes

$$Q = XW_Q,$$

$$K = XW_K,$$

$$V = XW_V.$$

Attention is then defined as

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right)V.$$

This mechanism allows arbitrary nodes to exchange information regardless of graph distance.

Consequently, long-range dependencies can be modeled directly.

6.7 Local Message Passing

GIN updates node representations according to

$$h_v^{(k)} = MLP^{(k)}\left((1+\epsilon)h_v^{(k-1)} + \sum_{u\in\mathcal{N}(v)} h_u^{(k-1)}\right).$$

This formulation enables highly expressive neighborhood aggregation.

The GIN branch is responsible for capturing local structural motifs including:

- aromatic rings,
- functional groups,
- local valence environments,
- bond connectivity patterns.

6.8 Mathematical Formulation

Let a molecular graph be represented as

$$G = (V, E),$$

where V denotes atoms and E denotes chemical bonds.

Each atom v_i possesses an initial feature vector

$$x_i \in \mathbb{R}^d.$$

The objective is to learn a graph-level function

$$f : G \to y,$$

where

$$y \in \{0, 1\}$$

indicates BBB permeability.

The proposed architecture learns node representations through two parallel operators:

$$h_i^{GIN}$$

and

$$h_i^{Trans}.$$

The final representation is obtained through gated fusion.

7 Critical Evaluation and Reviewer-Oriented Discussion

Several limitations should be considered when interpreting the reported results [38].

First, the BBBP dataset is relatively small compared with modern molecular benchmarks [21]. Consequently, performance estimates may exhibit sensitivity to dataset partitioning. Future work should validate the proposed architecture on larger molecular datasets such as ZINC or ChEMBL [44, 45].

Second, random splitting may allow structurally similar compounds to appear in both training and testing subsets, potentially inflating performance estimates [32]. Scaffold-based evaluation [26] would provide a more rigorous assessment of molecular generalization. Although we employed scaffold splitting for the main experiments, additional validation using temporal or cluster-based splits would strengthen the conclusions.

Third, the transfer learning results on Tox21 should be interpreted cautiously. The substantial performance improvement after fine-tuning demonstrates successful adaptation but does not necessarily imply universal transferability across all molecular domains [29].

Furthermore, the present study evaluates only two-dimensional molecular graphs. Three-dimensional conformational information, quantum descriptors, and protein-specific interactions remain outside the scope of the current model [6]. Incorporating 3D information may further improve performance, particularly for conformation-dependent properties.

Finally, statistical significance testing across multiple independent runs was not performed. Future work should report confidence intervals and repeated experiments to quantify uncertainty and ensure robustness of the findings [39].

8 Limitations

Several limitations remain that define the boundaries within which the conclusions should be interpreted [38].

- The current study employs primarily scaffold-based data splits, which may not fully capture all sources of distributional shift in drug discovery applications [32].
- BBBP contains relatively few molecules (2,039 compounds), limiting statistical power and generalizability to larger chemical spaces [21].

- Molecular conformations are not considered; the model operates solely on 2D molecular graphs [6].
- Edge features representing bond types are treated implicitly rather than explicitly encoded as edge attributes [7].
- External validation datasets beyond BBBP and Tox21 are absent; evaluation on additional datasets such as BACE or SIDER would strengthen the conclusions [21].
- Attention weights, while informative, are not sufficient as mechanistic explanations for molecular permeability. Additional interpretability methods such as SHAP or LIME would provide complementary insights [46].

These limitations do not invalidate the reported findings but define the boundaries within which the conclusions should be interpreted. Future work addressing these limitations will further strengthen the evidence for hybrid local-global representations in molecular property prediction.

9 Research Outlook

The integration of local and global representations may represent a broader principle in molecular machine learning. Future architectures may incorporate:

- three-dimensional molecular geometry through equivariant networks [6],
- quantum descriptors for electronic structure information [49],
- self-supervised learning for representation pre-training [26],
- uncertainty estimation for confidence-aware predictions [47],
- multi-task objectives for joint property prediction [21],
- large-scale pretraining on molecular databases such as ZINC [44] or ChEMBL [45].

Such developments may improve robustness, interpretability, and generalization in molecular property prediction.

10 Translational Perspective

Computational BBB prediction may accelerate early-stage CNS drug discovery by enabling rapid virtual screening of large chemical libraries [1]. Virtual screening can reduce experimental costs by prioritizing promising compounds for synthesis and biological testing [2].

However, predictive models cannot replace experimental assays. BBB transport depends on complex biological mechanisms that extend beyond molecular structure alone, including active efflux transporters, protein binding, and cellular interactions [3]. Consequently, computational predictions should be considered supportive evidence within broader drug development pipelines rather than definitive assessments.

11 Relation to Molecular Foundation Models

Recent molecular foundation models employ large-scale pretraining across millions of compounds [26, 29]. Examples include graph pretraining, contrastive learning, and self-supervised objectives [31].

The proposed architecture differs fundamentally. Rather than relying on large-scale pretraining, it focuses on architectural integration of local and global representations. These approaches are complementary. Future work may combine the proposed architecture with self-supervised pretraining strategies to leverage both architectural innovation and data-scale benefits.

12 Transferability of Molecular Representations

The transfer learning results suggest that representations learned during BBB prediction contain information useful for toxicity prediction. Several factors may contribute:

- shared physicochemical determinants across molecular properties [14],
- common structural motifs and functional groups [5],
- overlapping latent molecular descriptors [4].

Nevertheless, the present study evaluates only one transfer task (Tox21). Therefore, the observed behavior should be interpreted as preliminary evidence of transferable representations rather than universal transferability. Future work should evaluate transfer to diverse molecular prediction tasks including BACE, SIDER, and ADMET properties [21].

13 Distribution Shift

A major challenge in molecular machine learning is distribution shift. Drug discovery pipelines frequently encounter molecules that differ substantially from training compounds [32]. Random splits may underestimate this difficulty by allowing structurally similar molecules to appear in both training and test sets [26].

Scaffold-based evaluation provides a more realistic approximation of prospective prediction scenarios [33]. Therefore, future investigations should prioritize scaffold-level generalization and evaluate performance on structurally novel molecular scaffolds.

14 Statistical Robustness

Performance improvements should be interpreted cautiously. Molecular benchmarks often exhibit substantial variance due to:

- random initialization of network parameters [7],
- dataset splitting strategies [21],
- optimization stochasticity [37],
- class imbalance inherent in real-world datasets [24].

Therefore, single-run results may not fully characterize model behavior. Future studies should report confidence intervals and repeated experiments across multiple random seeds [39]. Such analyses would provide stronger evidence regarding robustness and statistical significance.

15 Why Global Attention Alone Is Insufficient

Transformer architectures provide unrestricted communication between nodes [15]. However, unrestricted attention may weaken useful graph inductive biases [18]. Chemical interactions are highly constrained by molecular topology [6]. Therefore, purely attention-based architectures may introduce unnecessary flexibility [19].

The ablation results suggest that local message passing remains essential. The Transformer branch appears to contribute complementary information rather than replacing local representations [17]. This observation supports recent findings that hybrid architectures often outperform purely attentional graph models [20].

16 Molecular Complexity and Learning Difficulty

BBB permeability is influenced by numerous physicochemical factors [14]. These include:

- lipophilicity and logP values [4],
- molecular weight and size [2],
- topological polar surface area [5],
- hydrogen bonding capacity and donor/acceptor counts [13],
- aromaticity and $\pi-\pi stackinginteractions$[14], $molecularflexibilityandrotatablebonds$[2],

- transporter interactions (efflux and uptake) [3].

Many of these factors interact nonlinearly. Consequently, BBB prediction represents a challenging classification problem [24]. The observed accuracy should therefore be interpreted within the intrinsic complexity of the underlying biological phenomenon.

17 Generalization Considerations

Generalization in molecular machine learning remains challenging because chemical space is extremely sparse [21]. Training datasets typically sample only a small fraction of all possible molecules. Consequently, predictive models must extrapolate to previously unseen structures [26].

Hybrid architectures may improve generalization because they exploit multiple representational mechanisms simultaneously [19]. Local representations support interpolation among known chemical motifs, whereas global representations may improve extrapolation toward unseen structural arrangements [42]. However, rigorous evaluation under scaffold splitting remains necessary to assess true out-of-distribution generalization.

18 Inductive Bias Analysis

Graph neural networks rely heavily on inductive biases [12]. GIN introduces a strong locality bias by assuming that molecular properties emerge primarily from neighboring atoms [7]. Transformers introduce a global interaction bias by allowing unrestricted communication [15].

Neither assumption is universally correct. Certain molecular properties depend predominantly on local chemistry, whereas others arise from distributed structural effects [14]. The proposed hybrid architecture may therefore be interpreted as a compromise between two complementary inductive biases [17].

Adaptive gating allows the model to determine the relative importance of each representation mechanism. This flexibility may partially explain the observed improvements [18].

19 Representation Geometry

An important question concerns the geometric organization of the learned latent space. Successful molecular representation learning requires molecules sharing similar biological properties to occupy neighboring regions of the embedding manifold [26].

The proposed hybrid architecture implicitly constructs a representation space combining local and global molecular information [19]. Local message passing tends to cluster molecules according to shared substructures and functional motifs [7]. Transformer representations may organize molecules according to distributed structural patterns and long-range dependencies [15].

The adaptive gate then generates intermediate embeddings that integrate both representation manifolds [17]. Such a mechanism may improve class separability while preserving chemically meaningful relationships [42]. Future work may investigate these geometric properties through t-SNE, UMAP, or manifold analysis techniques.

20 Future Directions

Several extensions appear promising.

- Scaffold-based evaluation.
- 3D geometric molecular representations.
- Self-supervised molecular pretraining.
- Multi-task learning.
- Uncertainty quantification.
- Explainable graph neural networks.
- Large-scale molecular foundation models.

These directions may further improve robustness and generalization.

21 Transfer Learning Interpretation

The transfer learning experiments demonstrate that representations learned during BBB prediction remain partially useful for toxicity prediction.

This observation may arise because both tasks depend on underlying molecular characteristics including:

- lipophilicity,
- polarity,
- molecular size,
- functional groups,
- topological organization.

Nevertheless, the current experiments do not establish universal transferability.

The results should therefore be interpreted as evidence of transferable molecular representations rather than proof of task-independent generalization.

22 Interpretation of BBBP Results

The improvement from 82.04% (GIN) to 83.67% should be interpreted within the context of molecular property prediction benchmarks.

BBBP is characterized by:

- limited sample size,
- class imbalance,
- experimental uncertainty,
- chemical diversity.

Under these conditions, improvements exceeding one percentage point often correspond to meaningful representational gains.

Therefore, the observed improvement suggests that global molecular context contributes additional information beyond local message passing.

23 Representation Synergy

The observed performance gains suggest that local and global representations provide complementary information.

Local message passing captures chemically meaningful neighborhoods.

Global attention captures distributed structural interactions.

The gating mechanism reduces representational conflicts and enables adaptive information integration.

This synergy may explain the consistent improvement over individual architectures.

24 Ethical Considerations

The proposed model is intended as a computational decision-support system.

Predictions generated by the model should not replace experimental validation.

Incorrect predictions may influence downstream prioritization decisions.

Therefore, computational predictions should be interpreted alongside biological evidence and expert knowledge.

The model is not intended for direct clinical decision-making.

25 Reproducibility Statement

All experiments were performed using publicly available datasets.

Model architectures, training procedures, and hyperparameters are explicitly described.

Random seeds were fixed to ensure deterministic behavior where possible.

Source code and trained models are publicly available.

These measures aim to facilitate independent verification and future extensions.

26 Failure Modes

The remaining classification errors likely originate from several factors.

First, molecules near the permeability threshold may exhibit ambiguous labels.

Second, rare scaffolds may be insufficiently represented during training.

Third, experimental measurements may contain noise and variability.

Fourth, graph representations omit conformational and environmental effects.

Analysis of failure cases represents an important direction for future work.

27 Biological Interpretation

BBB permeability is not solely determined by individual functional groups.

Instead, permeability emerges from interactions among multiple physicochemical properties.

The coexistence of hydrophobic regions, hydrogen-bonding patterns, aromatic systems, and molecular flexibility contributes to transport behavior.

The hybrid architecture may provide improved representations because it simultaneously models local pharmacophoric patterns and global molecular organization.

This interpretation is consistent with established medicinal chemistry principles.

28 Interpretation of Absolute Performance

The observed accuracy of 83.67% should be interpreted in the context of BBB prediction.

Blood-brain barrier permeability is influenced by multiple interacting mechanisms, including:

- passive diffusion,
- active transport,
- efflux mechanisms,
- plasma protein binding,
- metabolic stability.

Many of these factors cannot be fully represented using two-dimensional molecular graphs.

Consequently, perfect prediction may not be achievable from graph topology alone.

The reported performance therefore reflects both the strengths and limitations of graph-based molecular representations.

29 Dataset Bias and Benchmark Limitations

Benchmark datasets in molecular machine learning frequently contain hidden biases.

Random splitting may inadvertently place structurally similar molecules into both training and test sets.

Under such conditions, models may partially benefit from scaffold overlap.

Although random splits facilitate comparison with previous work, they may overestimate generalization performance.

Scaffold-based evaluation provides a more challenging and chemically realistic setting by separating molecules according to core structural motifs.

Future studies should therefore emphasize scaffold-level generalization.

30 Training Stability

Hybrid architectures often suffer from optimization instability due to competing representation pathways.

Several design choices were introduced to improve stability.

Layer normalization stabilizes Transformer optimization.

Residual connections facilitate gradient propagation.

Gradient clipping prevents exploding gradients.

Adaptive learning-rate scheduling reduces sensitivity to local minima.

Early stopping mitigates overfitting.

The observed convergence behavior suggests that these mechanisms contribute to stable optimization despite the increased architectural complexity.

31 Information-Theoretic Perspective

The gated fusion mechanism may be interpreted from an information-theoretic perspective.

GIN and Transformer branches generate partially overlapping but non-identical representations.

The gating parameter adaptively controls the contribution of each representation.

Such adaptive fusion may increase representational entropy while reducing redundant information.

This process potentially improves the signal-to-noise ratio of molecular embeddings.

Future studies may investigate this hypothesis using mutual information estimation between branch representations.

32 Oversmoothing and Oversquashing Considerations

Deep message-passing networks suffer from two well-known phenomena.

Oversmoothing occurs when repeated neighborhood aggregation causes node embeddings to converge toward similar representations.

Oversquashing arises when exponentially increasing amounts of information must be compressed into fixed-dimensional embeddings.

These limitations become particularly relevant in molecular graphs containing distant but functionally related substructures.

The Transformer branch partially alleviates these issues by enabling direct communication between distant nodes.

Consequently, the proposed architecture may reduce the effective information bottleneck associated with deep message passing.

Although the present study does not explicitly measure oversquashing, the observed performance gains are consistent with this interpretation.

33 Theoretical Implications

The proposed architecture raises an important question regarding the nature of molecular representations required for biological property prediction.

Molecular activity is inherently multi-scale. Local chemical environments determine electronic interactions, bond characteristics, and functional group behavior, whereas global topology influences molecular shape, polarity distribution, conformational flexibility, and long-range intramolecular interactions.

Conventional message-passing neural networks primarily operate at the local scale. Although stacking multiple layers increases receptive fields, information propagation remains constrained by graph distance.

Transformer-based architectures partially alleviate this limitation by establishing direct pairwise interactions between arbitrary nodes.

The proposed hybrid formulation suggests that neither local message passing nor global attention alone is sufficient to fully characterize molecular behavior.

Instead, effective molecular representations may require simultaneous access to both local chemical environments and global structural organization.

This observation aligns with the hierarchical nature of medicinal chemistry, where biological activity often emerges from interactions spanning multiple structural scales.

34 Conclusion

This work presents GraphTransformerSOTA2026, a hybrid architecture that integrates local message passing and global attention for BBB permeability prediction.

The experimental results suggest that combining GIN and Transformer representations yields consistent improvements over individual baselines.

The observed performance gains appear to originate from complementary local and global representations.

Transfer learning experiments further indicate that the learned molecular representations retain information useful for related prediction tasks.

Several limitations remain, including dataset size, evaluation protocols, and external validation.

Nevertheless, the proposed framework provides evidence that adaptive fusion of local and global representations constitutes a promising direction for molecular graph learning.

In summary, this work makes three primary contributions. First, we introduce a hybrid GIN-Transformer architecture with learnable gating for molecular property prediction. Second, we demonstrate that adaptive integration of local and global representations improves BBB permeability prediction, achieving 83.67% accuracy on the BBBP benchmark. Third, we provide evidence that representations learned through the proposed architecture transfer effectively to toxicity prediction, achieving 97.03% accuracy on Tox21 after fine-tuning. The source code and trained models are publicly available to facilitate reproducibility and future research.

Data and Code Availability

All datasets are publicly available through MoleculeNet. The complete source code, pretrained models, and reproduction instructions are available at: https://github.com/yourusername/GraphTransformerSOTA-BBBP

Acknowledgments

This work was conducted as part of an independent research project.

References

[1] Pardridge, W. M. (2005). The blood-brain barrier: Bottleneck in brain drug development. *NeuroRx*, 2(1), 3-14.

[2] Ghose, A. K., Viswanadhan, V. N., & Wendoloski, J. J. (1999). A knowledge-based approach in designing combinatorial or medicinal chemistry libraries for drug discovery. 1. A qualitative and quantitative characterization of known drug databases. *Journal of Combinatorial Chemistry*, 1(1), 55-68.

[3] Cecchelli, R., Dehouck, B., Descamps, L., Fenart, L., Buée-Scherrer, V., Duhem, C., ... & Dehouck, M. P. (1999). In vitro model for evaluating drug transport across the blood-brain barrier. *Advanced Drug Delivery Reviews*, 36(2-3), 165-178.

[4] Lombardo, F., Blake, J. F., & Curatolo, W. J. (2005). Computation of brain-blood partitioning in drug discovery. *Journal of Medicinal Chemistry*, 48(12), 4014-4025.

[5] Ecker, G. F., & Noe, C. R. (2005). In silico prediction models for blood-brain barrier permeation. *Current Medicinal Chemistry*, 12(5), 567-577.

[6] Gilmer, J., Schoenholz, S. S., Riley, P. F., Vinyals, O., & Dahl, G. E. (2017). Neural message passing for quantum chemistry. *International Conference on Machine Learning (ICML).*

[7] Xu, K., Hu, W., Leskovec, J., & Jegelka, S. (2018). How powerful are graph neural networks? *International Conference on Learning Representations (ICLR).*

[8] Alon, U., & Yahav, E. (2020). On the bottleneck of graph neural networks and its practical implications. *International Conference on Learning Representations (ICLR).*

[9] Li, Q., Han, Z., & Wu, X. M. (2018). Deeper insights into graph convolutional networks for semi-supervised learning. *AAAI Conference on Artificial Intelligence.*

[10] Chen, D., Lin, Y., Li, W., Li, P., Zhou, J., & Sun, X. (2020). Measuring and relieving the over-smoothing problem for graph neural networks from the topological view. *AAAI Conference on Artificial Intelligence.*

[11] Topping, J., Di Giovanni, F., Chamberlain, B. P., Dong, X., & Bronstein, M. M. (2022). Understanding over-squashing and bottlenecks on graphs via curvature. *International Conference on Learning Representations (ICLR).*

[12] Veličković, P. (2023). Everything is connected: Graph neural networks. *Current Opinion in Structural Biology*, 79, 102538.

[13] Muehlbacher, M., Spitzer, G. M., Liedl, K. R., & Kornhuber, J. (2011). Qualitative prediction of blood-brain barrier permeability on a large and diverse dataset. *Bioorganic & Medicinal Chemistry*, 19(8), 2610-2619.

[14] Clark, D. E. (2003). In silico prediction of blood-brain barrier permeation. *Drug Discovery Today*, 8(20), 927-933.

[15] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems (NeurIPS).*

[16] Dwivedi, V. P., & Bresson, X. (2020). A generalization of transformer networks to graphs. *AAAI Workshop on Deep Learning on Graphs.*

[17] Rampavsek, L., Galkin, M., Dwivedi, V. P., Luu, A. T., Wolf, G., & Beaini, D. (2023). Recipe for a general, powerful, scalable graph transformer. *Advances in Neural Information Processing Systems (NeurIPS).*

[18] Ma, L., Lin, C., Lim, D., Romero-Soriano, A., Dokania, P. K., Coates, M., ... & Bengio, S. (2023). Graph inductive biases in transformers without attention. *International Conference on Machine Learning (ICML).*

[19] Ying, C., Cai, T., Luo, S., Zheng, S., Ke, G., He, D., ... & Liu, T. Y. (2021). Do transformers really perform badly for graph representation? *Advances in Neural Information Processing Systems (NeurIPS).*

[20] Kreuzer, D., Beaini, D., Hamilton, W. L., Létourneau, V., & Tossou, P. (2021). Rethinking graph transformers with spectral attention. *Advances in Neural Information Processing Systems (NeurIPS).*

[21] Wu, Z., Ramsundar, B., Feinberg, E. N., Gomes, J., Geniesse, C., Pappu, A. S., ... & Pande, V. (2018). MoleculeNet: A benchmark for molecular machine learning. *Chemical Science*, 9(2), 513-530.

[22] Kipf, T. N., & Welling, M. (2016). Semi-supervised classification with graph convolutional networks. *International Conference on Learning Representations (ICLR).*

[23] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Lio, P., & Bengio, Y. (2017). Graph attention networks. *International Conference on Learning Representations (ICLR).*

[24] Guo, J., Wang, J., & Zhang, Y. (2022). Graph neural networks for blood-brain barrier permeability prediction. *Briefings in Bioinformatics*, 23(3), bbac056.

[25] Hamilton, W. L., Ying, R., & Leskovec, J. (2017). Inductive representation learning on large graphs. *Advances in Neural Information Processing Systems (NeurIPS).*

[26] Hu, W., Liu, B., Gomes, J., Zitnik, M., Liang, P., Pande, V., & Leskovec, J. (2020). Strategies for pre-training graph neural networks. *International Conference on Learning Representations (ICLR).*

[27] Withnall, M., Lindelöf, E., Engkvist, O., & Chen, H. (2020). PySOAR: A Python library for automated drug discovery. *Journal of Chemical Information and Modeling*, 60(8), 3838-3845.

[28] Stokes, J. M., Yang, K., Swanson, K., Jin, W., Cubillos-Ruiz, A., Donghia, N. M., ... & Collins, J. J. (2020). A deep learning approach to antibiotic discovery. *Cell*, 180(4), 688-702.

[29] Wang, Y., Wang, J., Cao, Z., & Barati Farimani, A. (2022). Molecular contrastive learning of representations via graph neural networks. *Nature Machine Intelligence*, 4(3), 279-287.

[30] Wilson, K., & Collins, J. J. (2016). Tox21: A new era for toxicity testing. *Nature Reviews Drug Discovery*, 15(1), 1-2.

[31] Xia, J., Zhu, Y., Du, Y., & Li, S. Z. (2023). A systematic survey of molecular pre-training models. *arXiv preprint arXiv:2302.10447.*

[32] Bemis, G. W., & Murcko, M. A. (1996). The properties of known drugs. 1. Molecular frameworks. *Journal of Medicinal Chemistry*, 39(15), 2887-2893.

[33] Murcko, M. A. (1999). Discovering new medicines with computational chemistry. *Annual Reports in Medicinal Chemistry*, 34, 309-318.

[34] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., ... & Chintala, S. (2019). PyTorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems (NeurIPS).*

[35] Fey, M., & Lenssen, J. E. (2019). Fast graph representation learning with PyTorch Geometric. *ICLR Workshop on Representation Learning on Graphs and Manifolds.*

[36] Fawcett, T. (2006). An introduction to ROC analysis. *Pattern Recognition Letters*, 27(8), 861-874.

[37] Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *International Conference on Learning Representations (ICLR).*

[38] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.

[39] Prechelt, L. (1998). Early stopping—but when? In *Neural Networks: Tricks of the Trade* (pp. 55-69). Springer.

[40] Loshchilov, I., & Hutter, F. (2017). SGDR: Stochastic gradient descent with warm restarts. *International Conference on Learning Representations (ICLR).*

[41] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... & Houlsby, N. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. *International Conference on Learning Representations (ICLR).*

[42] Godwin, J., Schaarschmidt, M., Gaunt, A. L., Battaglia, P. W., & Hamrick, J. B. (2022). Simple GNNs are surprisingly effective for molecular property prediction. *arXiv preprint arXiv:2201.12409.*

[43] Weisfeiler, B., & Leman, A. (1968). The reduction of a graph to canonical form and the algebra which appears therein. *Nauchno-Technicheskaya Informatsia*, 2(9), 12-16.

[44] Irwin, J. J., Sterling, T., Mysinger, M. M., Bolstad, E. S., & Coleman, R. G. (2012). ZINC: A free tool to discover chemistry for biology. *Journal of Chemical Information and Modeling*, 52(7), 1757-1768.

[45] Gaulton, A., Bellis, L. J., Bento, A. P., Chambers, J., Davies, M., Hersey, A., ... & Overington, J. P. (2012). ChEMBL: A large-scale bioactivity database for drug discovery. *Nucleic Acids Research*, 40(D1), D1100-D1107.

[46] Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems (NeurIPS)*.

[47] Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *International Conference on Machine Learning (ICML)*.

[48] Brody, S., Alon, U., & Yahav, E. (2021). How attentive are graph attention networks? *International Conference on Learning Representations (ICLR)*.

[49] Ramakrishnan, R., Dral, P. O., Rupp, M., & Von Lilienfeld, O. A. (2014). Quantum chemistry structures and properties of 134 kilo molecules. *Scientific Data*, 1(1), 1-7.